



Faculty of Science



Loop Parallelism I

Cosmin E. Oancea

`cosmin.oancea@diku.dk`

Department of Computer Science (DIKU)
University of Copenhagen

September 2023 PMPH Lecture Slides



Course Organization

W	HARDWARE		SOFTWARE	LAB/CUDA
1	Trends Vector Machine	←	List HOM (Map-Reduce)	Intro & Simple Map Programming
2	In Order Processor	→ ←	VLIW Instr Scheduling	Scan & Reduce
3	Cache Coherence		Loop Parallelism I	Sparse Vect Matrix Mult
4	Interconnection Networks		Case Studies & Optimizations	Transpose & Matrix Matrix Mult
5	Memory Consistency		Optimising Locality	Sorting & Profiling & Mem Optimizations
6	OoO, Spec Processor		Thread-Level Speculation	Project Work

Three narrative threads: the path to complex & good design:

- **Design Space** tradeoffs, constraints, common case, trends.
- **Reasoning**: from simple to complex, **Applying Concepts**.



Motivation

- + So far we reasoned about how to parallelize a known algorithm
- + using a clean, functional approach, e.g., flattening,
- + which provides work and depth guarantees,
 - but does **NOT** account for locality of reference.

Why do we have to look at imperative loops?



Motivation

- + So far we reasoned about how to parallelize a known algorithm
- + using a clean, functional approach, e.g., flattening,
- + which provides work and depth guarantees,
 - but does **NOT** account for locality of reference.

Why do we have to look at imperative loops?

- A lot of legacy sequential imperative code, C++/Java/Fortran.
- Need to parallelize the implementation of unknown algorithm,
- Need to optimize parallelism, e.g., locality of reference requires subscript analysis.



- 1 Direction-Vector Analysis
- 2 Optimizing Temporal Locality by Block Tiling
 - Brute-Force Nearest Neighbor Case Study
 - Matrix Multiplication Case Study
- 3 Coalesced Accesses: Matrix Transposition Case Study



Problem Statement

Three Loop Examples

DO i = 1, N	DO i = 2, N	DO i = 2, N
DO j = 1, N	DO j = 2, N	DO j = 1, N
A[j,i] = A[j,i] ..	A[j,i] = A[j-1,i-1]...	A[i,j] = A[i-1,j+1]...
ENDDO	B[j,i] = B[j-1,i]...	ENDDO
ENDDO	ENDDO ENDDO	ENDDO

Iterations are ordered *lexicographically*, i.e., in the order they occur in the sequential execution, e.g., $\vec{k} = (i=2, j=4) < \vec{l} = (i=3, j=3)$.

- Which of the three loop nests is amenable to parallelization?
- Loop interchange is one of the most simple and useful code transformations, e.g., used to enhance locality of reference, parallel-loop granularity, and even to “create” parallelism.
- In which loop nest is it safe to interchange the loops?



Definition of a Dependency

Load-Store Classification of Dependencies

True Dependency (RAW)

S1 X = ..

S2 .. = X

Anti Dependency (WAR)

S1 .. = X

S2 X = ..

Output dependency (WAW)

S1 X = ...

S2 X = ...

Def. Loop Dependence: There is a dependence from statement S_1 to S_2 in a loop nest *iff* \exists iterations \vec{k}, \vec{l} such that:



Definition of a Dependency

Load-Store Classification of Dependencies

True Dependency (RAW)		Anti Dependency (WAR)		Output dependency (WAW)	
S1	X = ..	S1	.. = X	S1	X = ...
S2	.. = X	S2	X = ..	S2	X = ...

Def. Loop Dependence: There is a dependence from statement $S1$ to $S2$ in a loop nest *iff* \exists iterations \vec{k}, \vec{l} such that:

1. $\vec{k} < \vec{l}$ or $\vec{k} = \vec{l}$ and \exists an execution path from statement $S1$ to statement $S2$ **such that:**
2. $S1$ accesses memory location M on iteration \vec{k} , and
3. $S2$ accesses memory location M on iteration \vec{l} , and
4. one of these accesses is a write.

We say that $S1$ is the source and $S2$ is the sink of the dependence, because $S1$ executes before $S2$ in the sequential program execution. Dependence depicted with an arrow pointing from source to sink.



Definition of a Dependency

Load-Store Classification of Dependencies

True Dependency (RAW)		Anti Dependency (WAR)		Output dependency (WAW)	
S1	X = ..	S1	.. = X	S1	X = ...
S2	.. = X	S2	X = ..	S2	X = ...

Def. Loop Dependence: There is a dependence from statement $S1$ to $S2$ in a loop nest *iff* \exists iterations \vec{k}, \vec{l} such that:

1. $\vec{k} < \vec{l}$ or $\vec{k} = \vec{l}$ and \exists an execution path from statement $S1$ to statement $S2$ **such that:**
2. $S1$ accesses memory location M on iteration \vec{k} , and
3. $S2$ accesses memory location M on iteration \vec{l} , and
4. one of these accesses is a write.

We say that $S1$ is the source and $S2$ is the sink of the dependence, because $S1$ executes before $S2$ in the sequential program execution. Dependence depicted with an arrow pointing from source to sink.

We are most interested in cross iteration dependencies, i.e., $\vec{k} < \vec{l}$. Intra iteration dependencies, i.e., $\vec{k} = \vec{l}$ are analysed for ILP.



Loop-Nest Dependencies

Lexicographic ordering, e.g., $\vec{k}=(i=2,j=4) < \vec{l}=(i=3,j=3)$.

Three Loop Examples

```
DO i = 1, N
  DO j = 1, N
    A[j,i] = A[j,i] ..
  ENDDO
ENDDO
```

```
DO i = 2, N
  DO j = 2, N
    A[j,i] = A[j-1,i-1]...
    B[j,i] = B[j-1,i]...
  ENDDO ENDDO
```

```
DO i = 2, N
  DO j = 1, N
    A[i,j] = A[i-1,j+1]...
  ENDDO
ENDDO
```



Loop-Nest Dependencies

Lexicographic ordering, e.g., $\vec{k} = (i=2, j=4) < \vec{l} = (i=3, j=3)$.

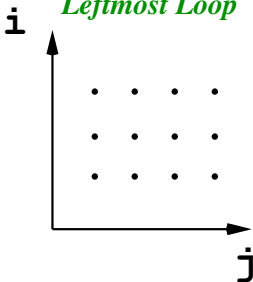
Three Loop Examples

```
DO i = 1, N
  DO j = 1, N
    A[j,i] = A[j,i] ..
  ENDDO
ENDDO
```

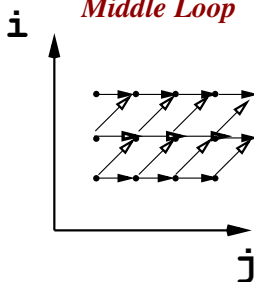
```
DO i = 2, N
  DO j = 2, N
    A[j,i] = A[j-1,i-1]...
    B[j,i] = B[j-1,i]...
  ENDDO
ENDDO
```

```
DO i = 2, N
  DO j = 1, N
    A[i,j] = A[i-1,j+1]...
  ENDDO
ENDDO
```

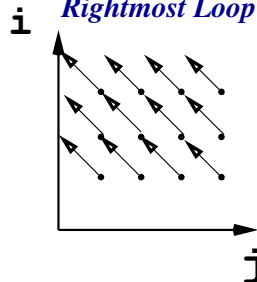
Leftmost Loop



Middle Loop



Rightmost Loop



How can I summarize this information?



Aggregate Dependencies via Direction Vectors

Write the Direction Vectors for Each Loop:

DO i = 1, N	DO i = 2, N	DO i = 2, N
DO j = 1, N	DO j = 2, N	DO j = 1, N
S1 A[j,i]=A[j,i]..	S1 A[j,i]=A[j-1,i]...	S1 A[i,j]=A[i-1,j+1]...
ENDDO	S2 B[j,i]=B[j-1,i-1]...	ENDDO
ENDDO	ENDDO ENDDO	ENDDO

Dependencies depicted via an edge *from* the stmt that executes first in the loop nest, i.e., *the source*, to the one that executes later, *the sink*.

Def. Dependence Direction: Assume \exists a dependence from $S1$ in iteration \vec{k} to $S2$ in \vec{l} ($\vec{k} \leq \vec{l}$). *Dependence-direction vector* $\vec{D}(\vec{k}, \vec{l})$:

1. $\vec{D}(\vec{k}, \vec{l})_m = "<"$ if $\vec{k}_m < \vec{l}_m$,
2. $\vec{D}(\vec{k}, \vec{l})_m = "="$ if $\vec{k}_m = \vec{l}_m$,
3. $\vec{D}(\vec{k}, \vec{l})_m = ">"$ if $\vec{k}_m > \vec{l}_m$.

If the source is a write and the sink a read then RAW dependency,
if the source is a read then WAR, if both are writes then WAW.



How to Compute the Direction Vectors?

- For any two statements $S1$ and $S2$ that may access the same array A (and one of the accesses is a write),
- in two symbolic iterations $I^1 \equiv (i_1^1, \dots, i_m^1)$ and $I^2 \equiv (i_1^2, \dots, i_m^2)$ (such that $I^1 < I^2$)
- on indices $A[e_1^1, \dots, e_n^1]$ and $A[e_1^2, \dots, e_n^2]$, respectively,
- then *the direction vectors may be derived* from the equations

$$\begin{cases} e_1^1 = e_1^2 \\ \dots \\ e_n^1 = e_n^2 \end{cases}$$

(The system of equations models the definition of a dependency: both accesses need to refer to the same memory location!)



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

```
DO i = 1, N
  DO j = 1, N
S1    A[j,i]=A[j,i]..
      ENDDO
  ENDDO
```



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

```

DO i = 1, N
  DO j = 1, N
S1    A[j,i]=A[j,i]..
      ENDDO
  ENDDO
For S1→S1:
  (j1,i1)=(j2,i2)
  i1 = i2 & j1 = j2

```

```

DO i = 2, N
  DO j = 2, N
S1    A[j,i]=A[j-1,i]...
S2    B[j,i]=B[j-1,i-1]...
      ENDDO
  ENDDO

```

Direction matrix:

S1→S1: [=,=]



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

```
DO i = 1, N
  DO j = 1, N
S1    A[j,i]=A[j,i]..
      ENDDO
```

```
ENDDO
```

```
For S1→S1:
```

```
(j1,i1)=(j2,i2)
```

```
i1 = i2 & j1 = j2
```

```
Direction matrix:
```

```
S1→S1: [=,=]
```

```
DO i = 2, N
  DO j = 2, N
S1    A[j,i]=A[j-1,i]...
S2    B[j,i]=B[j-1,i-1]...
```

```
ENDDO
```

```
ENDDO
```

```
S1→S1: (j1,i1)=(j2-1,i2)
```

```
i1 = i2 & j1 < j2
```

```
S2→S2: (j1,i1)=(j2-1,i2-1)
```

```
i1 < i2 & j1 < j2
```

```
S1→S1: [=,<]
```

```
S2→S2: [<,<]
```

```
DO i = 2, N
  DO j = 1, N
S1    A[i,j]=A[i-1,j+1]...
```

```
ENDDO
```

```
ENDDO
```

```
For S1→S1:
```



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

<pre> DO i = 1, N DO j = 1, N S1 A[j,i]=A[j,i].. ENDDO ENDDO For S1→S1: (j1,i1)=(j2,i2) i1 = i2 & j1 = j2 Direction matrix: S1→S1: [=,=]</pre>	<pre> DO i = 2, N DO j = 2, N S1 A[j,i]=A[j-1,i]... S2 B[j,i]=B[j-1,i-1]... ENDDO ENDDO S1→S1: (j1,i1)=(j2-1,i2) i1 = i2 & j1 < j2 S2→S2: (j1,i1)=(j2-1,i2-1) i1 < i2 & j1 < j2 S1→S1: [=,<] S2→S2: [<,<]</pre>	<pre> DO i = 2, N DO j = 1, N S1 A[i,j]=A[i-1,j+1]... ENDDO ENDDO For S1→S1: (i1,j1) = (i2-1,j2+1) i1 < i2 & j1 > j2 Direction matrix: S1→S1: [<,>]</pre>
---	---	--

Th. Parallelism: A loop in a loop nest is parallel *iff* all its directions are either = or there exists an outer loop whose corresp. direction is <.

A direction vector cannot have > as the first non-= symbol, as that would mean that I depend on something in the future.



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

```

DO i = 1, N
  DO j = 1, N
S1    A[j,i]=A[j,i]..
      ENDDO
    ENDDO

```

```

For S1→S1:
  (j1,i1)=(j2,i2)
  i1 = i2 & j1 = j2

```

Direction matrix:

S1→S1: [=,=]

```

DO i = 2, N
  DO j = 2, N
S1    A[j,i]=A[j-1,i]...
S2    B[j,i]=B[j-1,i-1]...
      ENDDO
    ENDDO

```

```

S1→S1: (j1,i1)=(j2-1,i2)
        i1 = i2 & j1 < j2
S2→S2: (j1,i1)=(j2-1,i2-1)
        i1 < i2 & j1 < j2

```

S1→S1: [=,<]

S2→S2: [<,<]

```

DO i = 2, N
  DO j = 1, N
S1    A[i,j]=A[i-1,j+1]...
      ENDDO
    ENDDO

```

```

For S1→S1:
  (i1,j1) = (i1-1,j2+1)
  i1 < i2 & j1 > j2

```

Direction matrix:

S1→S1: [<,>]

Th. Loop Interchange: A column permutation of the loops in a loop nest is legal *iff* permuting the direction matrix in the same way *does NOT* result in a > direction as the leftmost non-= direction in a row.



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

DO i = 1, N	DO i = 2, N	DO i = 2, N
DO j = 1, N	DO j = 2, N	DO j = 1, N
S1 A[j,i]=A[j,i]..	S1 A[j,i]=A[j-1,i]...	S1 A[i,j]=A[i-1,j+1]...
ENDDO	S2 B[j,i]=B[j-1,i-1]...	ENDDO
ENDDO	ENDDO ENDDO	ENDDO
For S1→S1: j1 = j2	For S1→S1: j1 = j2-1	For S1→S1: i1 = i2-1
i1 = i2	i1 = i2	j1 = j2+1
(i2,j2)-(i1,j1)=	(i2,j2)-(i1,j1)=[=,<]	(i2,j2)-(i1,j1)=[<,>]
[=,=]	For S2→S2: j1 = j2-1	
	i1 = i2-1	
	(i2,j2)-(i1,j1)=[<,<]	

Interchange is safe for the first and second nests, but not for the third!

e.g., $\begin{bmatrix} =, < \\ <, < \end{bmatrix} \rightarrow \begin{bmatrix} <, = \\ <, < \end{bmatrix}$ (for the second loop nest)



Parallelism and Loop Interchange

Direction Vectors/Matrix for Three Loops

DO i = 1, N	DO i = 2, N	DO i = 2, N
DO j = 1, N	DO j = 2, N	DO j = 1, N
S1 A[j,i]=A[j,i]..	S1 A[j,i]=A[j-1,i]...	S1 A[i,j]=A[i-1,j+1]...
ENDDO	S2 B[j,i]=B[j-1,i-1]...	ENDDO
ENDDO	ENDDO ENDDO	ENDDO
For S1→S1: j1 = j2	For S1→S1: j1 = j2-1	For S1→S1: i1 = i2-1
i1 = i2	i1 = i2	j1 = j2+1
(i2,j2)-(i1,j1)=	(i2,j2)-(i1,j1)=[=,<]	(i2,j2)-(i1,j1)=[<,>]
[=,=]	For S2→S2: j1 = j2-1	
	i1 = i2-1	
	(i2,j2)-(i1,j1)=[<,<]	

Interchange is safe for the first and second nests, but not for the third!

e.g., [=,<] → [<,<] (for the second loop nest)
 [<,<] [<,<]

After interchange, loop j of the second loop nest is parallel.

Corollary: A parallel loop can be always interchanged inwards.



Dependency Graph and Loop Distribution

Def. Dependency Graph: edges from the source of the dependency, i.e., early iteration, to the sink, i.e., later iteration.

Th. Loop Distribution: Statements that are in a dependence cycle remain in one (sequential) loop. The others are distributed to separate loops in graph order. A loop with no dependency cycle is parallel!

Vectorization Example: Remember Vector Machines?

```
DO i = 3, N
S1  A[i] = B[i-2] ...
S2  B[i] = B[i-1] ...
ENDDO

For S2→S1: i1 = i2-2, [<]
For S2→S2: i1 = i2-1, [<]
```



Dependency Graph and Loop Distribution

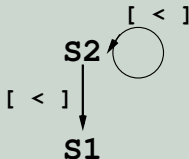
Def. Dependency Graph: edges from the source of the dependency, i.e., early iteration, to the sink, i.e., later iteration.

Th. Loop Distribution: Statements that are in a dependence cycle remain in one (sequential) loop. The others are distributed to separate loops in graph order. A loop with no dependency cycle is parallel!

Vectorization Example: Remember Vector Machines?

```
DO i = 3, N
S1  A[i] = B[i-2] ...
S2  B[i] = B[i-1] ...
ENDDO
```

```
For S2→S1: i1 = i2-2, [<]
For S2→S2: i1 = i2-1, [<]
```



```
DO i = 3, N
S2  B[i] = B[i-1] ...
ENDDO
```

```
DOALL i = 3, N
S1  A[i] = B[i-2] ...
ENDDOALL
```

Corollary: It is always legal to distribute a parallel loop;
but requires array expansion for local variables or if output
dependencies are present.



Loop Distribution May Require Array Expansion

```
float tmp;  
for(int i=2; i<N; i++) {  
    tmp = 2*B[i-2];  
    A[i] = tmp;  
    B[i] = tmp+B[i-1];  
}
```



Loop Distribution May Require Array Expansion

```
float tmp;  
for(int i=2; i<N; i++) {  
    tmp = 2*B[i-2];  
    A[i] = tmp;  
    B[i] = tmp+B[i-1];  
}  
  
float tmp[N];  
for(int i=2; i<N; i++) {  
    tmp[i] = 2*B[i-2];  
    B[i] = tmp[i]+B[i-1];  
}  
  
forall(int i=2; i<N; i++) {  
    A[i] = tmp[i];  
}
```

No matter where `tmp` is declared (inside or outside the loop) it needs to be expanded into an array in order to perform loop distribution.

If `tmp` is declared outside the loop then it requires **privatization**,



Loop Distribution May Require Array Expansion

```
float tmp;  
for(int i=2; i<N; i++) {  
    tmp = 2*B[i-2];  
    A[i] = tmp;  
    B[i] = tmp+B[i-1];  
}
```

```
float tmp[N];  
for(int i=2; i<N; i++) {  
    tmp[i] = 2*B[i-2];  
    B[i] = tmp[i]+B[i-1];  
}  
  
forall(int i=2; i<N; i++) {  
    A[i] = tmp[i];  
}
```

No matter where `tmp` is declared (inside or outside the loop) it needs to be expanded into an array in order to perform loop distribution.

If `tmp` is declared outside the loop then it requires **privatization**, because it actually causes frequent WAW dependencies. However its value is written before being used within the same iteration. Hence it is semantically equivalent to a locally declared variable, which will remove the output (WAW) dependency.

Distribution requires array expansion of the scalar `tmp`.



False Dependencies (WAR/WAW)

- **Cross-Iteration Anti Dependencies (WAR)** correspond to a read from the array as it was before the loop \Rightarrow can be eliminated by reading from a copy of the array.
- **Cross-Iteration WAW Dependencies (WAW):**
If they correspond to the case in which every **read** from a scalar or array location is covered by a **previous same-iteration write** \Rightarrow can be eliminated **privatization (renaming)**, which semantically moves the declaration of the variable (scalar or array) inside the loop.
- Direction-vectors reasoning is limited to relatively simple loop nests, e.g., difficult to reason about privatization in such a way.



Example: Eliminating WAR Dependencies

Anti Dependency (WAR) and Rewritten with Parallel Loops

```
float tmp = A[1];  
for (i=0; i<N-1; i++)  
S1  A[i] = A[i+1];  
    A[N-1] = tmp;  
  
//S1→S1: i1+1=i2, [<] WAR
```



Example: Eliminating WAR Dependencies

Anti Dependency (WAR) and Rewritten with Parallel Loops

```

float tmp = A[1];
for (i=0; i<N-1; i++)
S1  A[i] = A[i+1];
    A[N-1] = tmp;

//S1→S1: i1+1=i2, [<] WAR

// Solution: copy A into A'
// and use A' for the reads!
float Acopy[N];
#pragma omp parallel for
    for(int i=0; i<N; i++) {
        Acopy[i] = A[i];
    }
    tmp = A[1];
#pragma omp parallel for private(i)
    for (i=0; i<N-1; i++) {
        A[i] = Acopy[i+1];
    }
    A[N-1] = tmp;

```

For OpenMP code, run the g++ compiler with the -fopenmp flag. You may play with the degree of parallelism by setting the OMP_NUM_THREADS environment variable, e.g., to run on 4 cores use:

```
$ export OMP_NUM_THREADS=4.
```



Eliminating WAW Dependencies by Privatization

Assumption: array A is not used after the target loop.

```
int A[M];
for(i=0; i<N; i++){
    for(int j=0, j<M; j++)
        A[j] = (4*i+4*j) % M;
    for(int k=0; k<N; k++)
        X[i,k] = X[i,k-1] *
            A[A[(2*i+k)%M]%M];
}
```



Eliminating WAW Dependencies by Privatization

Assumption: array A is not used after the target loop.

// The write to A[j] causes multiple WAWs,

```
int A[M];
for(i=0; i<N; i++){
  for(int j=0, j<M; j++)
    A[j] = (4*i+4*j) % M;
  for(int k=0; k<N; k++)
    X[i,k] = X[i,k-1] *
      A[A[(2*i+k)%M]%M];
}
```



Eliminating WAW Dependencies by Privatization

Assumption: array A is not used after the target loop.

```

int A[M];
for(i=0; i<N; i++){
    for(int j=0, j<M; j++)
        A[j] = (4*i+4*j) % M;
    for(int k=0; k<N; k++)
        X[i,k] = X[i,k-1] *
            A[A[(2*i+k)%M]%M];
}

// The write to A[j] causes multiple WAWs,
// but A is fully written in the inner loop
#pragma omp parallel
{
    int A[M];
    #pragma omp for
    for(int i=0; i<N; i++){
        for(int j=0, j<M; j++)
            A[j] = (4*i+4*j) % M;
        for(int k=0; k<N; k++)
            X[i,k]=X[i,k-1] *
                A[A[(2*i+k)%M]%M];
    }
}

```

The declaration of `int A[M]` inside the OMP parallel region, will create an array A for every thread (and visible/accessible inside the thread context). Alternatively you may expand A with a new dimension of number-of-processor size, and use explicit indexing.



Reduction is Typically Easy To Recognize

If all the statements in which a scalar variable x appears are of the form $x \oplus = \text{exp}$, where x does not appear in exp and \oplus is associative (and commutative) then the cross-iteration RAWs on x can be resolved by:

- performing the updates $x \oplus = \text{exp}$ atomically, OR by:
 - privatizing x initialized with the neutral element,
 - computing the per-processor partial values of x ,
 - reducing the x s across processors and with the initial value.

A loop in which all dependencies are of the form above is called a **generalized reduction**; all properties of parallel loops are preserved.

```
// compilation requires g++ -fopenmp ...
float x = 6.0;
#pragma omp parallel for reduction(+:x) private(i,j)
for(i=1; i<N; i++) {
  for(j=1; j<N; j++) {
    if ( A[i,j] >= 2.0 )    x += 2*A[i,j-1];
    else if( A[i,j] > 0.0 ) x += A[i-1,j+1];
  }
  if ( i % (j+1) == 3 ) x += A[i,i];
}
```



Scan and Segmented Scan Are Difficult!

Compilers cannot recognize and parallelize even simple scans:

- they raise a cross-iteration true dependency (RAW),
- they appear in a multitude of forms,
- hence they are difficult to analyze.

// What kind of scans are these?

1. `A[0] = B[0];`
 `for(i=1; i<N; i++) {`
 `A[i] = A[i-1] + B[i];`
 `}`
2. `acc = 0;`
 `for(i=0; i<N; i++){`
 `acc = acc ^ i;`
 `A[i] = acc;`
 `}`
3. `for(j=0; j<M; j++)`
 `A[0,j] = B[0,j];`
 `for(i=1; i<N; i++) {`
 `for(j=0; j<M; j++)`
 `A[i,j] = A[i-1,j] + B[i,j];`
 `}`



Scan and Segmented Scan Are Difficult!

Compilers cannot recognize and parallelize even simple scans:

- they raise a cross-iteration true dependency (RAW),
- they appear in a multitude of forms,
- hence they are difficult to analyze.

// What kind of scans are these?

```
1. A[0] = B[0];
   for(i=1; i<N; i++) {
       A[i] = A[i-1] + B[i];
   }
```

```
2. acc = 0;
   for(i=0; i<N; i++){
       acc = acc ^ i;
       A[i] = acc;
   }
```

```
3. for(j=0; j<M; j++)
    A[0,j] = B[0,j];
   for(i=1; i<N; i++) {
       for(j=0; j<M; j++)
           A[i,j] = A[i-1,j] + B[i,j];
   }
```

```
1. let A = scan (+) 0 B
```

```
2. let A = scan (^) 0 (iota N)
```

```
3. let A = scan (\ a b -> map2 (+) a b)
      (replicate M 0.0) B
```

≡

```
let A = transpose <|
    map (scan (+) 0.0) <|
    transpose B
```



- 1 Direction-Vector Analysis
- 2 Optimizing Temporal Locality by Block Tiling
 - Brute-Force Nearest Neighbor Case Study
 - Matrix Multiplication Case Study
- 3 Coalesced Accesses: Matrix Transposition Case Study



Brute-Force Nearest Neighbor

For simplicity we assume a 1-dimensional space:

- N reference points;
- M queries, i.e., for each query we need to find its nearest neighbor w.r.t. the reference points.

```
DOALL i = 0, M-1, 1 // In parallel
  float query = queries[i];
  float nn_dist = Infinity;
  float nn_ind = -1;
  // Sequential over ref. points
  DO j = 0, N-1, 1
    ref_pt = points[j];
    float dist = fabs(query-ref_pt);
    IF dist < nn_dist
      nn_ind = j;
      nn_dist = dist;
    ENDDO
  ENDDO
  nns[i] = nn_ind;
ENDDO
```

Strategy:

- We execute the queries in parallel;
- We sequentialize the computation of each query; in principle it is a map-reduce, but we do not exploit it.
- Why?



Brute-Force Nearest Neighbor

For simplicity we assume a 1-dimensional space:

- N reference points;
- M queries, i.e., for each query we need to find its nearest neighbor w.r.t. the reference points.

```
DOALL i = 0, M-1, 1 // In parallel
  float query = queries[i];
  float nn_dist = Infinity;
  float nn_ind = -1;
  // Sequential over ref. points
  DO j = 0, N-1, 1
    ref_pt = points[j];
    float dist = fabs(query-ref_pt);
    IF dist < nn_dist
      nn_ind = j;
      nn_dist = dist;
    ENDIF
  ENDDO
  nns[i] = nn_ind;
ENDDO
```

Strategy:

- We execute the queries in parallel;
- We sequentialize the computation of each query; in principle it is a map-reduce, but we do not exploit it.
- Why?

Please note that the queries access the ref points in the same order, so there is potential for exploiting temporal locality.



Loop Strip Mining: A Simple Transformation

Strip mining breaks a loop into two-nested loops, as shown below:

```
// Original Code
DO i = 0, N, 1 // stride 1
  loop_body(i)
ENDDO
```

```
// After Strip Mining
DO ii = 0, N, T // stride T
  DO i = ii, MIN(ii+T-1,N), 1
    loop_body(i)
  ENDDO
ENDDO
```

When is it safe to strip-mine a loop?

Strip Mining is always safe because the observant student will surely notice that the transformed loop nest executes the same instructions the original loop and in the same order.



Strip Mining: Forming the CUDA Grid and Block

```
DOALL ii = 0, M-1, B // CUDA Grid
  DOALL i = ii, MIN(ii+B-1,M-1), 1 //Block
    // local thread id within Block
    int    loc_tid = i - ii;
    float query = queries[i];
    float nn_dist = Infinity;
    float nn_ind  = -1;
    // Sequential over ref. points
    DO j = 0, N-1, 1
      ref_pt = points[j];
      float dist = fabs(query-ref_pt);
      IF dist < nn_dist
        nn_ind  = j;
        nn_dist = dist;
      ENDIF
    ENDDO
    nns[i] = nn_ind;
  ENDDO // END CUDA BLOCK
ENDDO // END CUDA GRID
```

- Semantically, blockIdx.x would be ii/B,
- Semantically, threadIdx.x would be i - ii;
- Hardware caching will exploit some temporal locality across threads, but we can do better by using shared memory as a staging buffer!



Strip Mining: Blocking the Sequential Loop

```

DOALL ii = 0, M-1, B // CUDA Grid
  DOALL i = ii, MIN(ii+B-1,M-1), 1 // CUDA Block
    // local thread id within Block
    int    loc_tid = i - ii;
    float query = queries[i];
    float nn_dist = Infinity;
    float nn_ind  = -1;
    // Sequential over ref. points
    DO jj = 0, N-1, B
      // What slice of the points array is
      // used by the loop below? What to do?
      DO j = jj, MIN(jj+B-1, N-1), 1
        ref_pt = points[j];
        float dist = fabs(query-ref_pt);
        IF dist < nn_dist
          nn_ind  = j;
          nn_dist = dist;
        ENDIF
      ENDDO ENDDO
    nns[i] = nn_ind;
  ENDDO
ENDDO

```



Strip Mining: Blocking the Sequential Loop

```

DOALL ii = 0, M-1, B // CUDA Grid
  DOALL i = ii, MIN(ii+B-1,M-1), 1 // CUDA Block
    // local thread id within Block
    int loc_tid = i - ii;
    float query = queries[i];
    float nn_dist = Infinity;
    float nn_ind = -1;
    // Sequential over ref. points
    DO jj = 0, N-1, B
      // What slice of the points array is
      // used by the loop below? What to do?
      DO j = jj, MIN(jj+B-1, N-1), 1
        ref_pt = points[j];
        float dist = fabs(query-ref_pt);
        IF dist < nn_dist
          nn_ind = j;
          nn_dist = dist;
        ENDIF
      ENDDO ENDDO
    nns[i] = nn_ind;
  ENDDO
ENDDO

```

- Strip mine the sequential loop by block size B;
- Notice that all B threads of the inner block, access in the innermost loop the slice `points[jj: jj+B-1: 1]`, which has at most B consecutive elements.
- **Solution:** the B threads of the CUDA block collectively copy the B elements into a shared-memory buffer, and then read from there in the innermost loop (instead of accessing the points array which is held in global memory).
- **Essentially we use shared memory as an explicitly programmable cache to reduce the number of global-memory accesses by a factor of B.**



Shared Memory as an Explicitly-Programmed Cache

```

DOALL ii = 0, M-1, B // CUDA Grid
  DOALL i = ii, MIN(ii+B-1,M-1), 1 // CUDA Block
    // local thread id within Block
    int  loc_tid = i - ii;
    float query = queries[i];
    float nn_dist = Infinity;
    float nn_ind = -1;
    // Sequential over ref. points
    DO jj = 0, N-1, B
      float buffer[B];
      collectiveCopy(buffer, points, jj, B);
      // buffer[loc_tid] = points[jj+loc_tid];
      __syncthreads();
      DO j = jj, MIN(jj+B-1, N-1), 1
        ref_pt = buffer[j-jj]; //points[j]
        float dist = fabs(query-ref_pt);
        IF dist < nn_dist
          nn_ind = j;
          nn_dist = dist;
        ENDIF
      ENDDO __syncthreads(); ENDDO
      nns[i] = nn_ind;
    ENDDO
  ENDDO

```

- Essentially we use shared memory as an explicitly programmable cache to reduce the number of global-memory accesses by a factor of B .
- Each of the B threads copies one element from global to shared memory to fill in the buffer. Please note that the access to global memory is coalesced.
- Each thread then reads in the innermost loop B elements from the shared-memory buffer. The number of accesses to global memory is thus reduced with a factor of B !



Approximate CUDA Code

```
// CPU code
int B = 256;
int dimx = (M + B - 1) / B;
// dimx = ceil( ((float)M) / B );
dim3 block(B,1,1), grid(dimx,1,1);

unsigned long int elapsed;
struct timeval t_start,t_end,t_diff;
gettimeofday(&t_start, NULL);
    nearestNeighb<B><<<grid, block>>>
        (queries, points, nns, M, N);

gettimeofday(&t_end, NULL);
timeval_subtract(&t_diff,
                &t_end,&t_start);
elapsed=(t_diff.tv_sec*1e6 +
        t_diff.tv_usec);
double flops = 3.0 * M * N;
double gigaFlops=(flops*1.0e-3f) /
        elapsed;
```

```
template <int B> // KERNEL
__global__ void nearestNeighb( ... ) {
    __shared__ float shmem[B];
    int ii = blockIdx.x * B; //blockDim.x==B
    int tidx = threadIdx.x, i = ii+tidx;
    float nn_dist = INFINITY;
    int nn_ind = -1;
    float query = (i<M)? queries[i] : 0.0;
    for(int jj=0; jj<N; jj+=B) {
        shmem[tidx] = (jj+tidx<N) ?
            points[jj+tidx] : INFINITY;
        __syncthreads();
        for(int j=0; j<B; j++) {
            //if(jj+j < N && i < M) {
                float ref_pt = shmem[j];
                float dist = fabs(query - ref_pt);
                if(dist < nn_dist) {
                    nn_ind = jj + j;
                    nn_dist = dist;
                }
            }
        } __syncthreads();
    } if ( i < M ) nns[i] = nn_dist;
}
```



Matrix Multiplication: Loop Strip Mining

```
DOALL i = 0, M-1, 1    // Parallel
  DOALL j = 0, N-1, 1  // Parallel
    float tmp = 0.0
    DO k = 0, U-1, 1  // Reduction
      tmp += A[i,k]*B[k,j]
    ENDDO
    C[i,j] = tmp;
  ENDDO
ENDDO
```

←Matrix Multiplication. Matrices:

- input A has M rows and U columns
- input B has U rows and N columns
- result C has M rows and N columns

Loops of indices i and j are parallel
(can be proved by direction vectors).

Accesses to A and B invariant to loops i and j \Rightarrow Block Tiling to optimize locality of reference!



Matrix Multiplication: Loop Strip Mining

```
DOALL i = 0, M-1, 1    // Parallel
  DOALL j = 0, N-1, 1  // Parallel
    float tmp = 0.0
    DO k = 0, U-1, 1  // Reduction
      tmp += A[i,k]*B[k,j]
    ENDDO
    C[i,j] = tmp;
  ENDDO
ENDDO
```

← Matrix Multiplication. Matrices:

- input A has M rows and U columns
- input B has U rows and N columns
- result C has M rows and N columns

Loops of indices i and j are parallel
(can be proved by direction vectors).

Accesses to A and B invariant to loops i and j \Rightarrow Block Tiling to optimize locality of reference!

First step: Strip Mining, always safe since the transformed loop executes the same instructions in the same order as the original loop:

```
DO i = 0, N-1, 1  // stride 1
  loop_body(i)
ENDDO
```

```
DO ii = 0, N-1, T  // stride T
  DO i = ii, MIN(ii+T-1, N-1), 1
    loop_body(i)
  ENDDO
ENDDO
```



Matrix Multiplication: Loop Interchange

After strip mining all loops with a tile of size T:

```
DOALL ii = 0, M-1, T
  DOALL i = ii, MIN(ii+T-1,M-1), 1    // loop
    DOALL jj = 1, N, T                // interchange. Why Safe?
      DOALL j = jj, MIN(jj+T-1,N-1), 1
        float tmp = 0.0
        DO kk = 0, U-1, T
          DO k = kk, MIN(kk+T-1,U-1), 1
            tmp += A[i,k]*B[k,j]
          ENDDO ENDDO
        C[i,j] = tmp;
      ENDDO ENDDO ENDDO ENDDO
```



Matrix Multiplication: Loop Interchange

After strip mining all loops with a tile of size T:

```
DOALL ii = 0, M-1, T
  DOALL i = ii, MIN(ii+T-1,M-1), 1    // loop
    DOALL jj = 1, N, T                // interchange. Why Safe?
      DOALL j = jj, MIN(jj+T-1,N-1), 1
        float tmp = 0.0
        DO kk = 0, U-1, T
          DO k = kk, MIN(kk+T-1,U-1), 1
            tmp += A[i,k]*B[k,j]
          ENDDO ENDDO
        C[i,j] = tmp;
      ENDDO ENDDO ENDDO ENDDO
```

The second step is to apply loop interchange between the loops of indices *i* and *jj*. This is safe because loop *i* is parallel, hence it can always be interchanged inwards!



Matrix Multiplication: Summarizing Read Subscripts

After loop interchange we have a grid shape, as in CUDA:

```
DOALL ii = 0, M-1, T           // grid.y
  DOALL jj = 0, N-1, T         // grid.x
    DOALL i = ii, MIN(ii+T-1,M-1), 1 // block.y
      DOALL j = jj, MIN(jj+T-1,N-1), 1 // block.x
        float tmp = 0.0
        DO kk = 0, U-1, T
          DO k = kk, MIN(kk+T-1,U-1), 1
            tmp += A[i,k]*B[k,j]
          ENDDO
        ENDDO
      ENDDO
    ENDDO ENDDO ENDDO ENDDO
```

The third step is to summarize the subscripts of A and B read inside the loop of index k, for fixed ii, jj and kk (x:y denotes [x...y]):



Matrix Multiplication: Summarizing Read Subscripts

After loop interchange we have a grid shape, as in CUDA:

```
DOALL ii = 0, M-1, T           // grid.y
  DOALL jj = 0, N-1, T         // grid.x
    DOALL i = ii, MIN(ii+T-1,M-1), 1 // block.y
      DOALL j = jj, MIN(jj+T-1,N-1), 1 // block.x
        float tmp = 0.0
        DO kk = 0, U-1, T
          DO k = kk, MIN(kk+T-1,U-1), 1
            tmp += A[i,k]*B[k,j]
          ENDDO
        ENDDO
      ENDDO
    ENDDO
  ENDDO
ENDDO ENDDO ENDDO ENDDO
```

The third step is to summarize the subscripts of A and B read inside the loop of index k, for fixed ii, jj and kk (x:y denotes [x..y]):

- A subscripts [ii : MIN(ii+T-1,M), kk : MIN(kk+T-1,U)]
- B subscripts [kk : MIN(kk+T-1,U), jj : MIN(jj+T-1,N)]
- Summaries have size at most T^2 & independent on i, j, and k \Rightarrow CUDA-block threads cooperatively copy-in data to shared mem!



Block Tiled Matrix Multiplication CUDA Kernel

Shared memory padded with zeros to remove the branch from loop k!

```
DOALL ii = 0, M-1, T    // grid.y
  DOALL jj = 0, N-1, T  // grid.x
    DOALL i = ii, MIN(ii+T-1,M-1), 1
      DOALL j = jj, MIN(jj+T-1,N-1), 1
        float tmp = 0.0
        DO kk = 0, U-1, T
          //we would like to copy
          //to shared memory here
          //& use it inside loop k
          DO k = kk, MIN(kk+T-1,U-1), 1
            tmp += A[i,k]*B[k,j]
          ENDDO
        ENDDO
      ENDDO
    ENDDO ENDDO ENDDO ENDDO
```



Block Tiled Matrix Multiplication CUDA Kernel

Shared memory padded with zeros to remove the branch from loop k!

```

DOALL ii = 0, M-1, T    // grid.y
DOALL jj = 0, N-1, T    // grid.x
DOALL i = ii, MIN(ii+T-1,M-1), 1
DOALL j = jj, MIN(jj+T-1,N-1), 1
    float tmp = 0.0
    DO kk = 0, U-1, T
        //we would like to copy
        //to shared memory here
        //& use it inside loop k
        DO k = kk, MIN(kk+T-1,U-1), 1
            tmp += A[i,k]*B[k,j]
        ENDDO
    ENDDO
    C[i,j] = tmp;
ENDDO ENDDO ENDDO ENDDO

__global__ void matMultTiledKer( ... ) {
    __shared__ T Ash[T][T], Bsh[T][T];
    int ii = blockIdx.y * T; //blockDim.x==T
    int jj = blockIdx.x * T; //blockDim.y==T
    int tidy = threadIdx.y, i = tidy+ii;
    int tidx = threadIdx.x, j = tidx+jj;
    float tmp = 0.0;

    for(int kk=0; kk<U; kk+=T) {
        Ash[tidy][tidx] = (i<M && kk+tidx<U) ?
            A[i*U + (kk+tidx)] : 0.0;
        Bsh[tidy][tidx] = (j<N && kk_tidx<U) ?
            B[(kk+tidy)*N + j] : 0.0;
        __syncthreads();
        for(int k=0; k<T; k++) {
            tmp += Ash[tidy][k] * Bsh[k][tidx]
        } __syncthreads();
    } if (i<M && j<N) C[i*N + j] = tmp;
}

```

A global memory access amortized by (T-1) shared memory accesses.



Measuring GFlops Performance

Sequential matrix multiplication $\sim 2 \times M \times N \times U$ floating point operations. What is the GFlops performance of our implementation?

```
// CPU code
int dimy = ceil( ((float)M) / T );
int dimx = ceil( ((float)N) / T );
dim3 block(T,T,1), grid(dimx,dimy,1);

unsigned long int elapsed;
struct timeval t_start,t_end,t_diff;
gettimeofday(&t_start, NULL);

// ignoring generic shared mem problems
matMultTiledKer<T><<<grid, block>>>
    (d_A, d_B, d_C, U, M, N);

gettimeofday(&t_end, NULL);
timeval_subtract(&t_diff,
                &t_end,&t_start);
elapsed=(t_diff.tv_sec*1e6 +
        t_diff.tv_usec);
double flops = 2.0 * M * N * U;
double gigaFlops=(flops*1.0e-3f) /
    elapsed;

template <int T> // KERNEL
__global__ void matMultTiledKer( ... ) {
    __shared__ float Ash[T][T], Bsh[T][T];
    int ii = blockIdx.y * T; //blockDim.x==T
    int jj = blockIdx.x * T; //blockDim.y==T
    int tidy = threadIdx.y, i = tidy+ii;
    int tidx = threadIdx.x, j = tidx+jj;
    float tmp = 0.0;

    for(int kk=0; kk<U; kk+=T) {
        Ash[tidy][tidx] = (i<M && kk+tidx<U) ?
            A[i*U + (kk+tidx)] : 0.0;
        Bsh[tidy][tidx] = (j<N && kk_tidy<U) ?
            B[(kk+tidy)*N + j] : 0.0;
        __syncthreads();
        for(int k=0; k<T; k++) {
            tmp += Ash[tidy][k] * Bsh[k][tidx];
        } __syncthreads();
    } if (i<M && j<N) C[i*N + j] = tmp;
}
```



- 1 Direction-Vector Analysis
- 2 Optimizing Temporal Locality by Block Tiling
 - Brute-Force Nearest Neighbor Case Study
 - Matrix Multiplication Case Study
- 3 Coalesced Accesses: Matrix Transposition Case Study



Matrix Transposition: Motivation

```
// Non-Coalesced Memory Access
// Transposition to coalesce it ⇒
DOALL i = 0 to N-1 // parallel
  float accum = 0
  DO j = 0, 63 // sequential
    float tmpA = A[i, j]
    accum = accum*accum + tmpA*tmpA
    B[i,j] = accum
  ENDDO
ENDDO

A' = transpose(A)
DOALL i = 0 to N-1 // parallel
  float accum = 0
  DO j = 0, 63 // sequential
    float tmpA = A'[j, i]
    accum = accum*accum + tmpA*tmpA
    B'[j, i] = accum
  ENDDO
B = transpose(B')
```

The transformed program performs about three-times more accesses to global memory than the original.

But exhibits only coalesced accesses!

Coalesced Access: a (half) warp accesses in a SIMD instruction consecutive memory (word) locations.



Transposition: Strip Mining, Interchange & Kernel

```
//Both loops are parallel
//Strip mining & interchange⇒
for(i = 0; i < rowsA; i++) {
    for(j = 0; j < colsA; j++) {
        trA[j*rowsA+i] = A[i*colsA+j];
    } }

for(ii=0; ii<rowsA; ii+=T) {
    for(jj=0; jj<colsA; jj+=T) {
        for(i=ii; i<min(ii+T,rowsA); i++) {
            for(j=jj; j<min(jj+T,colsA); j++) {
                trA[j*rowsA+i] = A[i*colsA+j];
            } } } } }
```



Transposition: Strip Mining, Interchange & Kernel

```
//Both loops are parallel
//Strip mining & interchange⇒
for(i = 0; i < rowsA; i++) {
    for(j = 0; j < colsA; j++) {
        trA[j*rowsA+i] = A[i*colsA+j];
    } }
```

```
for(ii=0; ii<rowsA; ii+=T) {
    for(jj=0; jj<colsA; jj+=T) {
        for(i=ii; i<min(ii+T,rowsA); i++) {
            for(j=jj; j<min(jj+T,colsA); j++) {
                trA[j*rowsA+i] = A[i*colsA+j];
            } } } }
```

```
__global__ void matTranspose(
    float* A, float* trA,
    int rowsA, int colsA ) {
    __shared__ float tile[T][T+1];
    int tidx = threadIdx.x;
    int tidy = threadIdx.y;
    int j = blockIdx.x*T + tidx;
    int i = blockIdx.y*T + tidy;
    if( j < colsA && i < rowsA )
        tile[tidy][tidx] = A[i*colsA+j];
    __syncthreads();
    i = blockIdx.y*T + threadIdx.x;
    j = blockIdx.x*T + threadIdx.y;
    if( j < colsA && i < rowsA )
        trA[j*rowsA+i] = tile[tidx][tidy];
}
```

- Trick is to write the element of the symmetric thread in the same block.
- What is the problem?



Transposition: Strip Mining, Interchange & Kernel

```
//Both loops are parallel
//Strip mining & interchange⇒
for(i = 0; i < rowsA; i++) {
    for(j = 0; j < colsA; j++) {
        trA[j*rowsA+i] = A[i*colsA+j];
    } }
```

```
for(ii=0; ii<rowsA; ii+=T) {
    for(jj=0; jj<colsA; jj+=T) {
        for(i=ii; i<min(ii+T,rowsA); i++) {
            for(j=jj; j<min(jj+T,colsA); j++) {
                trA[j*rowsA+i] = A[i*colsA+j];
            } } } }
```

```
__global__ void matTranspose(
    float* A, float* trA,
    int rowsA, int colsA ) {
    __shared__ float tile[T][T+1];
    int tidX = threadIdx.x;
    int tidY = threadIdx.y;
    int j = blockIdx.x*T + tidX;
    int i = blockIdx.y*T + tidY;
    if( j < colsA && i < rowsA )
        tile[tidY][tidX] = A[i*colsA+j];
    __syncthreads();
    i = blockIdx.y*T + threadIdx.x;
    j = blockIdx.x*T + threadIdx.y;
    if( j < colsA && i < rowsA )
        trA[j*rowsA+i] = tile[tidX][tidY];
}
```

- Trick is to write the element of the symmetric thread in the same block.
- What is the problem?
- Number of shared memory banks typically 16 or 32.
- T is also either 16 or 32 ⇒
- 16 consecutive threads will read the same memory bank at the same time.
- Solution: `tile[T][T+1]`;



Lessons Learned So Far

- Tiled transposition is about $2\times$ faster than the naive version,
- but the motivating example runs much faster than that when transposition coalesces accesses to arrays A and B. **Why?**



Lessons Learned So Far

- Tiled transposition is about $2\times$ faster than the naive version,
- but the motivating example runs much faster than that when transposition coalesces accesses to arrays A and B. **Why?**
- **Better to eliminate rather than hide latency.** Impact of hardware multi-threading limited by the amount of available resources!



Constant (Read-Only) Memory in CUDA

- 64KB of `__constant__` memory on device (global/slow), cached in each multiprocessor, e.g., 8KB (fast).
- May reduce the required memory bandwidth:
 - if found in cache, then no extra traffic,
 - if a (half) warp accesses the same location and misses in cache \Rightarrow only one request is sent and the result is broadcast back to all,
 - **serialized** accesses if a warp of threads read different locations!
 - latency can range from one to hundreds of cycles.
- Best Use: when an entire block accesses the same location in the same SIMD instruction: even on a miss, the first warp brings the data in cache @ minimal traffic, the rest find it in cache.

```
// C in __constant__ memory: Good!
DO i = 1, N, 1 // grid
  DO j = 1, M, 1 // block(s)
    A[i,j] = A[i,j] % C[i]
  ENDDO
ENDDO

// C in __constant__ memory: Bad!
DO i = 1, N, 1 // grid
  DO j = 1, M, 1 // block(s)
    A[i,j] = A[i,j] % C[j]
  ENDDO
// Either global memory or loop interchange
```

